

The linear regression model, and model assumptions

Econometrics for minor Finance, Lecture 3

Tunga Kantacı, Fall 2025

An empirical challenge: Educational outcomes

Can we improve student outcomes?

What are these policy tools?

Changing class size? Improving teacher quality? Investment in school?

An empirical challenge: Educational outcomes: Policies

Can smaller classes yield to higher test scores? What are the mechanisms?

Classroom management: Fewer students can mean fewer disruptions, leading to a better learning environment.

Teacher workload: With fewer students, teachers may be less overwhelmed and better able to prepare and assess.

Individual attention: Smaller classes may allow teachers to give more personalized instruction and feedback. Students might feel more engaged and less overlooked.

An empirical challenge: Educational outcomes: Data

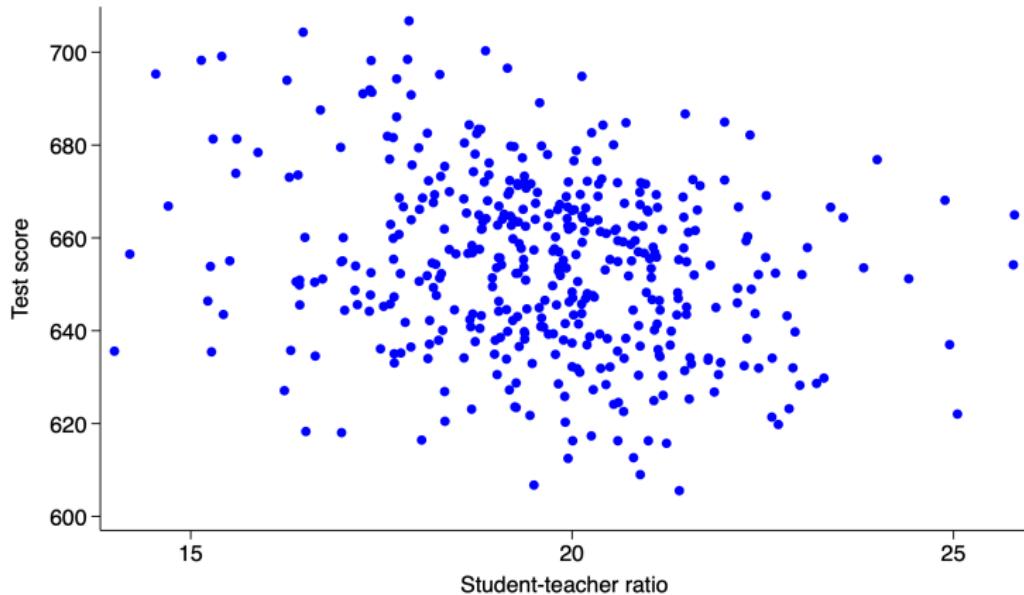
The California Standardized Testing and Reporting (STAR) dataset contains data on test performance and school characteristics. The data are from districts in California collected in 1999 by the California Department of Education. **Test scores** are the average of the reading and math scores on a standardized test administered to 5th grade students. The **student-teacher ratio** is the number of students divided by the number of teachers working full-time in a district. The data is analyzed in Kruger and Whitmore (Economic Journal, 2001).

An empirical challenge: Educational outcomes: Data: Summary statistics

```
. summarize tstscr str
```

Variable	Obs	Mean	Std. dev.	Min	Max
tstscr	420	654.1565	19.05335	605.55	706.75
str	420	19.64043	1.891812	14	25.8


An empirical challenge: Educational outcomes: Data: Scatter plot of test scores against student-teacher ratio



An empirical challenge: Educational outcomes: Model

Our aim would be to formally analyze how test scores vary with the student-teacher ratio. For this, we need a **model**, but not just any model. It must be one that fits the data well and captures the underlying relationship we care about.

An empirical challenge: Model and fitted values



The linear regression model: Model

An **observational unit** is the entity on which measurements are taken, such as a student, worker, firm, or country. The unit is random. It becomes non-random once we complete data sampling.

The linear regression model: Model

Our model is

$$y_i = \beta_0 \cdot 1_i + \beta_1 x_i + u_i$$

for each of n observational units.

y_i : outcome: e.g., test score

β_0 : intercept term, unknown so needs to be estimated

1_i : a constant equal to 1, used to include the intercept term

β_1 : slope coefficient, unknown so needs to be estimated

x_i : predictor: e.g., class size.

u_i : error term, to capture factors not observed in the data,
unobserved but we can approximate it

The linear regression model: Model

y_i also called dependent variable, or response variable

x_i also called independent variable, regressor, or explanatory variable

These names reflect convention, they do **not** imply causation.
Regression describes conditional relationships, and sometimes causal ones.

The linear regression model: Model

The model

$$y_i = \beta_0 \cdot 1_i + \beta_1 x_i + u_i$$

is a **population model**. It provides the conceptual recipe for how the data we observe are generated. That is why we sometimes call it the **data generating process in the population**: it assumes how the outcome variable y is produced from the explanatory variable x together with unobserved influences captured by the error term u . The **population parameters** β_0 and β_1 summarize the systematic part of this relationship in the population. Randomness enters through the sampling process that determines which units from the population are included in the dataset.

The linear regression model: Model: Notation note

The model for observational unit i is

$$y_i = \beta_0 \cdot 1_i + \beta_1 x_i + u_i$$

As we typically have sample data on n observational units, we can stack the observations to obtain a compact model for n units

$$y = \beta_0 \cdot 1 + \beta_1 x + u$$


Moreover, although the intercept is technically multiplied by a vector of ones, we suppress it in notation and write:

$$y = \beta_0 + \beta_1 x + u$$

The linear regression model: Model: Assumptions

Like any other model, a regression **model** abstracts away complexity by reducing reality to a functional form. Models are approximations of the reality.

The linear regression model: Model: Assumptions



The linear regression model: Model: Assumptions

"All models are wrong, but some are useful."
— George E. P. Box

The linear regression model: Model: Assumptions

What all this implies is that we are making **assumptions** under which any regression model yields interpretable, credible, and sometimes causal insights about the world. If these assumptions fail, the model is not a good approximation of reality. The point is not perfection, but usefulness under certain conditions. Regression is useful not because it is true, but because under the right assumptions it provides good approximations of relationships.

The linear regression model: Model: Assumptions: Linearity

Linearity. The model is linear in the parameters.

The linear regression model: Model: Assumptions: Linearity

The model

$$y = \beta_0 + \beta_1 x^2 + u$$

is

- linear in the parameters
- nonlinear in the regressor
- linear in the squared regressor

What makes a regression model linear is that it is linear in the **parameters**.

The linear regression model: Model: Assumptions: Linearity: Intuition

The model

$$y = x^\beta + u$$

is nonlinear in the parameter. This is not the linear regression model. Linearity means parameters enter additively and proportionally. If the true relationship looks more curved or complex, as it is here, this straight, proportional form is not a good assumption.

The linear regression model: Model: Assumptions: No perfect collinearity

No perfect collinearity. If there are multiple explanatory variables, they must be linearly independent. In other words, no column vector corresponding to one variable can be expressed as a linear combination of the column vectors of the other variables.

The linear regression model: Model: Assumptions: No perfect collinearity

If we include several explanatory variables in a regression model, each variable must contribute unique information. This means that no variable can be exactly predicted from the others. If one variable is just the sum or multiple of other variables, then it does not add anything new. When this happens, the regression model cannot separate the individual effects of the variables, and the estimator we want to use, to estimate the unknown parameters of the model, breaks down because it cannot distinguish which variable is responsible for changes in the outcome.

The linear regression model: Model: Assumptions: No perfect collinearity: Example

Consider the regression model

$$wage = \beta_0 \cdot 1 + \beta_1 d^{female} + \beta_2 d^{male} + u$$

where

$$d_i^{female} = \begin{cases} 1 & \text{if } i = female \\ 0 & \text{if } i = male \end{cases}$$

$$d_i^{male} = \begin{cases} 0 & \text{if } i = female \\ 1 & \text{if } i = male \end{cases}$$

Sum of the values in each **row** of d^{female} and d^{male} is equal to the value in that **row** of 1. Hence, one value can be perfectly predicted from other values. This is **perfect collinearity**.

The linear regression model: Model: Assumptions: No perfect collinearity: Example

$$\begin{bmatrix} 1 & d^{female} & d^{male} \end{bmatrix} = \begin{bmatrix} 1 & 1 & 0 \\ 1 & 1 & 0 \\ 1 & 1 & 0 \\ 1 & 0 & 1 \\ 1 & 0 & 1 \\ 1 & 0 & 1 \end{bmatrix}$$

The linear regression model: Model: Assumptions: No perfect collinearity: Intuition

Think of explanatory variables as different **ingredients** in a recipe. If each ingredient is distinct, like flour, sugar, butter, you can tell what role each plays. But if one ingredient is just a mix of the others, say, flour + sugar, then it does not add new information, it is redundant. Regression needs each variable to be distinct so it can measure its separate effect.

The linear regression model: Model: Assumptions: Exogeneity

Exogeneity.

$$E[u | x] = 0$$

This condition says that, once we account for x in the model, the remaining error, u , is just random noise. It does not systematically lean up or down depending on x . In other words: given any value of x , the model is right *on average*. This assumption is also sometimes called the zero conditional mean assumption.

The linear regression model: Model: Assumptions: Exogeneity: The theoretical implication

This assumption has two implications. First, the law of iterated expectations requires that

$$E[u] = E_X [E[u | x]]$$

Since

$$E[u | x] = 0$$

we have

$$E[u] = 0$$

The linear regression model: Model: Assumptions: Exogeneity: The theoretical implication

Second, the law of iterated expectations requires that

$$E[ux] = E_X [E[ux | x]] = E_X [x E[u | x]]$$

Since

$$E[u | x] = 0$$

we have

$$E[ux] = 0$$

The linear regression model: Model: Assumptions: Exogeneity: The Theoretical implication

Now, from See Math refresher B in Wooldridge, we know that

$$\text{Cov}[u, x] = E[ux] - E[u]E[x]$$

Since

$$E[ux] = 0$$

and

$$E[u] = 0$$

we have

$$\text{Cov}[u, x] = 0$$

So the assumption implies that u is not correlated with x .

The linear regression model: Model: Assumptions: Exogeneity: Intuition

Suppose we want to predict exam scores from study hours. Suppose students who study more also have hidden advantages, like natural ability. Since they are hidden, they show up in the error term. This makes the error correlated with study hours. The result is that the regression model mixes hidden ability with effort. The results it that the slope is biased.

The linear regression model: Model: Assumptions: Homoskedasticity

In Greek, **homo** means same, and **skedasis** means dispersion.
Hence, homoskedasticity means same dispersion. Constant variance.

The linear regression model: Model: Assumptions: Homoskedasticity

Homoskedasticity.

$$\text{Var}(u | x) = \sigma^2$$

This condition says that the error term u has the same variance at all values of x . It means that, regardless of the level of x , the variability of u is constant.

The linear regression model: Model: Assumptions: Homoskedasticity: Intuition

Suppose we want to predict wages from education. Individuals with little education face limited job options, so their wages vary less. The more educated can access a wider variety of jobs paying different wages, so their wages vary more. This means the variance of u depends on x . The result is that the regression misjudges its own precision: when errors vary with education, it becomes harder to get a precise prediction of how wages vary with education.

The linear regression model: Model: Assumptions: Normality of errors

The errors are normally distributed. u has the mean and variance given by the exogeneity and homoskedasticity assumptions, and has a normal distribution:

$$u | x \sim N[0, \sigma^2]$$

We will use this assumption if n is small but drop it if n is large.
This is a technical assumption. We will get back to this.

Sampling: Assumption: Random sampling

Random sampling. The data we collect

$$\{(x_i, y_i) : i = 1, 2, \dots, n\}$$

is a random sample of the population that follows the population model. It says that each unit in the population has equal probability of selection, and selections are independent: observations are independently and identically distributed.

Sampling: Assumption: Random sampling: Intuition

Suppose we want to know about the average education in the population. If only sample university students, the sample no longer represents the population. Then the model predictions will be systematically biased.

Sampling: Assumption: Random sampling: It does not guarantee representativeness

- Randomness does not eliminate **sampling variation**. Any single random sample may be unrepresentative, we can be **unlucky**. Representativeness is **probabilistic**, not absolute.
- Random sampling ensures the sample is **expected** to represent the population. **Across repeated samples**, the **sampling distribution** of statistics, e.g., means, proportions, converges to the population values.
- Randomness protects against **systematic bias**, but not against chance variation in one particular sample.

Sampling: Assumption: Random sampling: It does not guarantee representativeness: Example

A single random sample can be unrepresentative, but the distribution of a sample statistic **across repeated samples** centers on the true population value:

- In a 50/50 male-female population, one random sample of size $n = 10$ might yield 7 males and 3 females ($\hat{p} = 0.7$), which is not representative of $p = 0.5$.
- If we repeat this process hundreds of times, the **sampling distribution** of \hat{p} clusters around the true proportion $p = 0.5$.
- Random sampling guarantees representativeness **in expectation**, not in every single draw.

Sampling: Assumption: Random sampling: It does not guarantee representativeness: Example

